

Optimization and related uses of autodiff: Illustrations

Changcheng Li and John C. Nash

2018/7/1

Introduction

autodiff is an **R** package to perform automatic differentiation of **R** functions by calling the automatic differentiation tools in the **Julia** programming language. The document **FandR** ([??ref](#)) describes how to install **autodiff**.

Here we will illustrate how **autodiff** can be used to provide gradients or other derivative information for use in optimization problems for which **R** is being used to attempt solutions.

Problem setup

Most methods for optimization require

- an objective function that is to be minimized or maximized. Because the minimum of $f(x)$ is the maximum of $-f(x)$, we will only talk of minimizing functions. Though the mathematics of this are trivial, the care and attention to avoid errors when translating a maximization problem to one involving minimization require serious effort and continuous checking.
- a starting set of values for the parameters to be optimized (and perhaps also for any exogenous data and/or fixed parameters)

We need to load the **autodiff** package and then initiate it. NOTE: This is quite slow the first time it is run. WORSE: It is always treated as a “first time” when called in knitr during the processing of a vignette from an Rmd-type file.

```
library(autodiff)
```

```
##
## Attaching package: 'autodiff'
## The following object is masked from 'package:stats':
##
##      deriv
ad_setup()
```

```
## Julia version 0.6.3 at location /usr/local/bin will be used.
## Loading setup script for JuliaCall...
## Finish loading setup script for JuliaCall.
```

And we can use package **numDeriv** to compare with **autodiff**.

```
require(numDeriv)
```

```
## Loading required package: numDeriv
##
## Attaching package: 'numDeriv'
## The following objects are masked from 'package:autodiff':
##
##      grad, hessian, jacobian
```

Test problem – ViaRes

This is simply to test how we can get the gradient of a function that is defined as the sum of squares of residuals, BUT the residuals are computed in a subsidiary function that must be called.

At July 2, 2018, this gives an error that stops knitr, so evaluation is turned off in the following example.

```
require(autodiffr)
ad_setup() # to ensure it is established

ores <- function(x){
  x # Function will be the parameters. ofn is sum of squares
}

ofn0 <- function(x){ # original ofn
  res <- ores(x) # returns a vector of residual values
  val <- as.numeric(crossprod(res)) # as.numeric because crossprod is a matrix
  val
}

ofn <- function(x){ # But autodiffr does not understand crossprod()
  res <- ores(x) # returns a vector of residual values
  val <- sum(res*res) # NOT crossprod()
  val
}
```

Note that this works with eval=TRUE, but Chebyquad still failing.

```
## Now try to generate the gradient function
ogr <- autodiffr::grad(ofn)

# print(ogr) # this will be more or less meaningless link to Julia function
x0 <- c(1,2,3)
print(ofn(x0)) # should be 14
```

```
## [1] 14
```

```
print(ofn0(x0)) # should be 14
```

```
## [1] 14
```

```
ogr0<-ogr(x0) # should be 2, 4, 6
ogr0
```

```
## [1] 2 4 6
```

Test problem – Chebyquad

This problem was given prominence in the optimization literature by Fletcher (1965).

First let us define our Chebyquad function. Note that this is for the **vector** x. This version is expressed as the sum of squares of a vector of function values, which provide a nonlinear least squares problem. Note that `crossprod()` may cause difficulties as it is not written in **R**.

```
require(autodiffr)
ad_setup()
cyq.f <- function(x) {
  rv<-cyq.res(x)
```

```

f <- sum(rv*rv)
}

cyq.res <- function (x) {
# Fletcher's chebyquad function m = n -- residuals
n<-length(x)
res<-rep(0,n) # initialize
for (i in 1:n) { #loop over resids
  rr<-0.0
  for (k in 1:n) {
    z7<-1.0
    z2<-2.0*x[k]-1.0
    z8<-z2
    j<-1
    while (j<i) {
      z6<-z7
      z7<-z8
      z8<-2*z2*z7-z6 # recurrence to compute Chebyshev polynomial
      j<-j+1
    } # end recurrence loop
    rr<-rr+z8
  } # end loop on k
  rr<-rr/n
  if (2*trunc(i/2) == i) { rr <- rr + 1.0/(i*i - 1) }
  res[i]<-rr
} # end loop on i
res
}

```

Let us choose a single value for the number of parameters, and for illustration use $n = 4$.

```

## cyq.setup
n <- 4
lower<-rep(-10.0, n)
upper<-rep(10.0, n)
x<-1:n
x<-x/(n+1.0) # Initial value suggested by Fletcher

```

For safety, let us check the function and a numerical approximation to the gradient.

```

require(numDeriv)
cat("Initial parameters:")

## Initial parameters:
print(x)

## [1] 0.2 0.4 0.6 0.8
cat("Initial value of the function is ",cyq.f(x),"\n")

## Initial value of the function is 0.07118393
gn <- numDeriv::grad(cyq.f, x) # using numDeriv
cat("Approximation to gradient at initial point:")

## Approximation to gradient at initial point:

```

```
print(gn)
```

```
## [1] 0.6170624 0.1882112 -0.1882112 -0.6170624
```

Using a modular approach to the problem, first specifying it via **residuals** and computing the function as a sum of squares, we can also generate the gradient.

```
# Ref: Fletcher, R. (1965) Function minimization without calculating derivatives -- a review,  
#       Computer J., 8, 33-41.
```

```
# Note we do not have all components here e.g., .jsd, .h
```

```
cyq.jac<- function (x) {  
  # Chebyquad Jacobian matrix  
  n<-length(x)  
  cj<-matrix(0.0, n, n)  
  for (i in 1:n) { # loop over rows  
    for (k in 1:n) { # loop over columns (parameters)  
      z5<-0.0  
      cj[i,k]<-2.0  
      z8<-2.0*x[k]-1.0  
      z2<-z8  
      z7<-1.0  
      j<- 1  
      while (j<i) { # recurrence loop  
        z4<-z5  
        z5<-cj[i,k]  
        cj[i,k]<-4.0*z8+2.0*z2*z5-z4  
        z6<-z7  
        z7<-z8  
        z8<-2.0*z2*z7-z6  
        j<- j+1  
      } # end recurrence loop  
      cj[i,k]<-cj[i,k]/n  
    } # end loop on k  
  } # end loop on i  
  cj  
}
```

```
cyq.g <- function (x) {  
  cj<-cyq.jac(x)  
  rv<-cyq.res(x)  
  gg<- 2.0 * as.vector(rv %*% cj)  
}
```

```
# check gradient function cyq.g
```

```
gajn <- cyq.g(x)  
print(gajn)
```

```
## [1] 0.6170624 0.1882112 -0.1882112 -0.6170624
```

We can now try to see if autodiffr matches this gradient. However, the following code gives an error in

Julia.

```
# Do not evaluate, as this fails
cyq.ag <- autodiffr::grad(cyq.f)
gaag <- cyq.ag(x)
print(gaag)
```

As a workaround, we can get the Chebyquad function from the package `funconstrain`. The `funconstrain` offering does NOT require a call to the residuals, but has a single level **R** function.

```
require(funconstrain)
```

```
## Loading required package: funconstrain
```

```
cat("funconstrain loaded\n")
```

```
## funconstrain loaded
```

```
cheb <- chebyquad() # Seem to need the brackets or doesn't return pieces
print(str(cheb))
```

```
## List of 4
## $ fn:function (par)
## $ gr:function (par)
## $ fg:function (par)
## $ x0:function (n = 50)
## NULL
```

```
cyq2.f <- cheb$fn
## Note that funconstrain offers the starting value
## x0b <- cheb$x0(n=4) # Need the size of the vector
## x0b
## cyq2.f(x0b)
## same as
print(cyq2.f(x))
```

```
## [1] 0.07118393
```

```
## Try the gradient
```

```
cyq2.ag <- autodiffr::grad(cyq2.f) # Need autodiffr:: specified for knitr
## print(cyq2.g)
cat("Gradient at x")
```

```
## Gradient at x
```

```
g2ag <- cyq2.ag(x)
print(g2ag)
```

```
## [1] 0.6170624 0.1882112 -0.1882112 -0.6170624
```

```
require(microbenchmark)
```

```
## Loading required package: microbenchmark
```

```
cat("cyq.f timing:\n")
```

```
## cyq.f timing:
```

```
tcyq.f <- microbenchmark(cyq.f(x))
tcyq.f
```

```
## Unit: microseconds
```

```
##      expr    min      lq    mean median      uq    max neval
## cyq.f(x) 16.507 16.7565 17.25711 16.918 17.0485 42.808   100
```

```
cat("cyq2.f timing:\n")
```

```
## cyq2.f timing:
```

```
tcyq2.f <- microbenchmark(cyq2.f(x))
tcyq2.f
```

```
## Unit: microseconds
```

```
##      expr    min      lq    mean median      uq    max neval
## cyq2.f(x) 6.226 6.456 6.83078 6.538 6.7245 27.36   100
```

```
tcyq.g <- microbenchmark(cyq.g(x))
tcyq.g
```

```
## Unit: microseconds
```

```
##      expr    min      lq    mean median      uq    max neval
## cyq.g(x) 44.349 45.106 82.29287 45.433 45.88 3706.073   100
```

```
tcyq2.ag <- microbenchmark(cyq2.ag(x), unit="us" )
# microseconds
tcyq2.ag
```

```
## Unit: microseconds
```

```
##      expr    min      lq    mean median      uq    max neval
## cyq2.ag(x) 26283.3 26548.4 27559.72 26728.18 27156.77 44006.25   100
```

Chebyquad Questions:

- How fast is the calculation? Try different ways and produce a table.
- Why is JN's cyq.ag NOT working? Does it have something to do with calling cyq.res to compute cyq.f? NO. crossprod() seems to be somehow invoked even though not explicit.

Test problem – Hobbs weed infestation

This nonlinear estimation problem was brought to one of the authors (JN) in the mid-1970s (See Nash (1979)). It has just 12 data points and asks for the estimation of a 3-parameter logistic growth curve. The present example does not provide for scaling.

```
hobbs.f<- function(x){ # Hobbs weeds problem -- function
  if (abs(12*x[3]) > 500) { # check computability
    fbad<-.Machine$double.xmax
    return(fbad)
  }
  res<-hobbs.res(x)
  f<-sum(res*res)
}
```

```
hobbs.res<-function(x){ # Hobbs weeds problem -- residual
# This variant uses looping
  if(length(x) != 3) stop("hobbs.res -- parameter vector n!=3")
  y<-c(5.308, 7.24, 9.638, 12.866, 17.069, 23.192, 31.443, 38.558, 50.156, 62.948,
       75.995, 91.972)
  t<-1:12
```

```

    if(abs(12*x[3])>50) {
      res<-rep(Inf,12)
    } else {
      res<-x[1]/(1+x[2]*exp(-x[3]*t)) - y
    }
  }

hobbs.jac<-function(x){ # Jacobian of Hobbs weeds problem
  jj<-matrix(0.0, 12, 3)
  t<-1:12
  yy<-exp(-x[3]*t)
  zz<-1.0/(1+x[2]*yy)
  jj[t,1] <- zz
  jj[t,2] <- -x[1]*zz*zz*yy
  jj[t,3] <- x[1]*zz*zz*yy*x[2]*t
  jjret <- jj
  attr(jjret,"gradient") <- jj
  return(jjret)
}

hobbs.g<-function(x){ # gradient of Hobbs weeds problem
  # NOT EFFICIENT TO CALL AGAIN
  jj<-hobbs.jac(x)
  res<-hobbs.res(x)
  gg<-as.vector(2.*t(jj) %*% res)
  return(gg)
}

hobbs.rsd<-function(x) { # Jacobian second derivative
  rsd<-array(0.0, c(12,3,3))
  t<-1:12
  yy<-exp(-x[3]*t)
  zz<-1.0/(1+x[2]*yy)
  rsd[t,1,1]<- 0.0
  rsd[t,2,1]<- -yy*zz*zz
  rsd[t,1,2]<- -yy*zz*zz
  rsd[t,2,2]<- 2.0*x[1]*yy*yy*zz*zz*zz
  rsd[t,3,1]<- t*x[2]*yy*zz*zz
  rsd[t,1,3]<- t*x[2]*yy*zz*zz
  rsd[t,3,2]<- t*x[1]*yy*zz*zz*(1-2*x[2]*yy*zz)
  rsd[t,2,3]<- t*x[1]*yy*zz*zz*(1-2*x[2]*yy*zz)
  ## rsd[t,3,3]<- 2*t*t*x[1]*x[2]*x[2]*yy*yy*zz*zz*zz
  rsd[t,3,3]<- -t*t*x[1]*x[2]*yy*zz*zz*(1-2*yy*zz*x[2])
  return(rsd)
}

hobbs.h <- function(x) { ## compute Hessian
  # cat("Hessian not yet available\n")
  # return(NULL)
  H<-matrix(0,3,3)
  res<-hobbs.res(x)
  jj<-hobbs.jac(x)

```

```

    rsd<-hobbs.rsd(x)
##    H<-2.0*(t(res) %*% rsd + t(jj) %*% jj)
    for (j in 1:3) {
      for (k in 1:3) {
        for (i in 1:12) {
          H[j,k]<-H[j,k]+res[i]*rsd[i,j,k]
        }
      }
    }
    H<-2*(H + t(jj) %*% jj)
    return(H)
}

x0good <- c(200, 50, 0.3)
x0bad <- c(1,1,1)
f0good <- hobbs.f(x0good)
cat("Sum of squares at the GOOD starting point:",f0good,"\n")

## Sum of squares at the GOOD starting point: 158.2324
f0bad <- hobbs.f(x0bad)
cat("Sum of squares at the BAD starting point:",f0bad,"\n")

## Sum of squares at the BAD starting point: 23520.58
res0good <- hobbs.res(x0good)
## Residuals -- good starting point
res0good

## [1] -0.05050236 -0.20779506 -0.26086802 -0.41247546 -0.61690702
## [6] -1.60525418 -3.36423901 -2.43016453 -4.28734016 -5.63079076
## [11] -5.67542775 -7.44779537
res0bad <- hobbs.res(x0bad)
## Residuals -- bad starting point
res0bad

## [1] -4.576941 -6.359203 -8.685426 -11.883986 -16.075693 -22.194473
## [7] -30.443911 -37.558335 -49.156123 -61.948045 -74.995017 -90.972006
require(autodiffr)
ad_setup()
hobbs.ag <- autodiffr::grad(hobbs.f)
hobbsag0good <- hobbs.ag(x0good)
## Gradient by AD -- good starting point
hobbsag0good

## [1] -17.8438 48.2491 -24559.8187
## Compare hand coded function
hobbsggood <- hobbs.g(x0good)
hobbsggood

## [1] -17.8438 48.2491 -24559.8187
## Gradient by AD -- bad starting point
hobbsag0bad <- hobbs.ag(x0bad)
hobbsag0bad

```



```
## [1] -824.042084    4.764888 -11.025384
```

```
## Compare hand coded function
```

```
hobbsgbad <- hobbs.g(x0bad)
```

```
hobbsgbad
```

```
## [1] -824.042084    4.764888 -11.025384
```

```
## Interestingly, the magnitude of gradient elements greater for "good"
```

```
hobbs.aj <- autodiffr::jacobian(hobbs.res)
```

```
## Gradient by AD -- good starting point
```

```
hobbsaj0good <- hobbs.aj(x0good)
```

```
hobbsaj0good
```

```
##           [,1]      [,2]      [,3]
```

```
## [1,] 0.02628749 -0.1023858  5.119291
```

```
## [2,] 0.03516102 -0.1356989 13.569891
```

```
## [3,] 0.04688566 -0.1787496 26.812437
```

```
## [4,] 0.06226762 -0.2335615 46.712293
```

```
## [5,] 0.08226046 -0.3019747 75.493681
```

```
## [6,] 0.10793373 -0.3851362 115.540847
```

```
## [7,] 0.14039380 -0.4827335 168.956738
```

```
## [8,] 0.18063918 -0.5920347 236.813864
```

```
## [9,] 0.22934330 -0.7069798 318.140911
```

```
## [10,] 0.28658605 -0.8178179 408.908969
```

```
## [11,] 0.35159786 -0.9119072 501.548971
```

```
## [12,] 0.42262102 -0.9760500 585.629985
```

```
## Compare hand coded function
```

```
hobbsjgood <- hobbs.jac(x0good)
```

```
hobbsjgood
```

```
##           [,1]      [,2]      [,3]
```

```
## [1,] 0.02628749 -0.1023858  5.119291
```

```
## [2,] 0.03516102 -0.1356989 13.569891
```

```
## [3,] 0.04688566 -0.1787496 26.812437
```

```
## [4,] 0.06226762 -0.2335615 46.712293
```

```
## [5,] 0.08226046 -0.3019747 75.493681
```

```
## [6,] 0.10793373 -0.3851362 115.540847
```

```
## [7,] 0.14039380 -0.4827335 168.956738
```

```
## [8,] 0.18063918 -0.5920347 236.813864
```

```
## [9,] 0.22934330 -0.7069798 318.140911
```

```
## [10,] 0.28658605 -0.8178179 408.908969
```

```
## [11,] 0.35159786 -0.9119072 501.548971
```

```
## [12,] 0.42262102 -0.9760500 585.629985
```

```
## attr(,"gradient")
```

```
##           [,1]      [,2]      [,3]
```

```
## [1,] 0.02628749 -0.1023858  5.119291
```

```
## [2,] 0.03516102 -0.1356989 13.569891
```

```
## [3,] 0.04688566 -0.1787496 26.812437
```

```
## [4,] 0.06226762 -0.2335615 46.712293
```

```
## [5,] 0.08226046 -0.3019747 75.493681
```

```
## [6,] 0.10793373 -0.3851362 115.540847
```

```
## [7,] 0.14039380 -0.4827335 168.956738
```

```

## [8,] 0.18063918 -0.5920347 236.813864
## [9,] 0.22934330 -0.7069798 318.140911
## [10,] 0.28658605 -0.8178179 408.908969
## [11,] 0.35159786 -0.9119072 501.548971
## [12,] 0.42262102 -0.9760500 585.629985

## Gradient by AD -- bad starting point
hobbsaj0bad <- hobbs.aj(x0bad)
hobbsaj0bad

##           [,1]           [,2]           [,3]
## [1,] 0.7310586 -1.966119e-01 1.966119e-01
## [2,] 0.8807971 -1.049936e-01 2.099872e-01
## [3,] 0.9525741 -4.517666e-02 1.355300e-01
## [4,] 0.9820138 -1.766271e-02 7.065082e-02
## [5,] 0.9933071 -6.648057e-03 3.324028e-02
## [6,] 0.9975274 -2.466509e-03 1.479906e-02
## [7,] 0.9990889 -9.102212e-04 6.371548e-03
## [8,] 0.9996646 -3.352377e-04 2.681901e-03
## [9,] 0.9998766 -1.233793e-04 1.110414e-03
## [10,] 0.9999546 -4.539581e-05 4.539581e-04
## [11,] 0.9999833 -1.670114e-05 1.837126e-04
## [12,] 0.9999939 -6.144137e-06 7.372964e-05

## Compare hand coded function
hobbsjbad <- hobbs.jac(x0bad)
hobbsjbad

##           [,1]           [,2]           [,3]
## [1,] 0.7310586 -1.966119e-01 1.966119e-01
## [2,] 0.8807971 -1.049936e-01 2.099872e-01
## [3,] 0.9525741 -4.517666e-02 1.355300e-01
## [4,] 0.9820138 -1.766271e-02 7.065082e-02
## [5,] 0.9933071 -6.648057e-03 3.324028e-02
## [6,] 0.9975274 -2.466509e-03 1.479906e-02
## [7,] 0.9990889 -9.102212e-04 6.371548e-03
## [8,] 0.9996646 -3.352377e-04 2.681901e-03
## [9,] 0.9998766 -1.233793e-04 1.110414e-03
## [10,] 0.9999546 -4.539581e-05 4.539581e-04
## [11,] 0.9999833 -1.670114e-05 1.837126e-04
## [12,] 0.9999939 -6.144137e-06 7.372964e-05
## attr(,"gradient")
##           [,1]           [,2]           [,3]
## [1,] 0.7310586 -1.966119e-01 1.966119e-01
## [2,] 0.8807971 -1.049936e-01 2.099872e-01
## [3,] 0.9525741 -4.517666e-02 1.355300e-01
## [4,] 0.9820138 -1.766271e-02 7.065082e-02
## [5,] 0.9933071 -6.648057e-03 3.324028e-02
## [6,] 0.9975274 -2.466509e-03 1.479906e-02
## [7,] 0.9990889 -9.102212e-04 6.371548e-03
## [8,] 0.9996646 -3.352377e-04 2.681901e-03
## [9,] 0.9998766 -1.233793e-04 1.110414e-03
## [10,] 0.9999546 -4.539581e-05 4.539581e-04
## [11,] 0.9999833 -1.670114e-05 1.837126e-04
## [12,] 0.9999939 -6.144137e-06 7.372964e-05

```

Now let us try this in a solution of nonlinear least squares.

WARNING: Because of some compatibility issues with other **R** software, the jacobian must be available in the “gradient” attribute returned by the jacobian function. The purpose of this is to allow the function `nlsr::nlfb` to have the same name for the residual and jacobian function. This is used in generating a symbolic jacobian function in `nlsr::nlxb`. However, it can catch unwary users (including us!).

```
# try in a function
require(nlsr)

## Loading required package: nlsr
## manual
smggood <- nlfb(x0good, hobbs.res, hobbs.jac, trace=TRUE)

## no weights
## lower:[1] -Inf -Inf -Inf
## upper:[1] Inf Inf Inf
## Start:lamda: 1e-04  SS= 158.2324  at  = 200  = 50  = 0.3  1 / 0
## <<lamda: 4e-05  SS= 2.61779  at  = 194.3011  = 48.56497  = 0.313994  2 / 1
## <<lamda: 1.6e-05  SS= 2.587325  at  = 196.0825  = 49.07632  = 0.313616  3 / 2
## <<lamda: 6.4e-06  SS= 2.587277  at  = 196.1839  = 49.09134  = 0.313571  4 / 3
## <<lamda: 2.56e-06  SS= 2.587277  at  = 196.1862  = 49.09164  = 0.3135697  5 / 4
## <<lamda: 1.024e-06  SS= 2.587277  at  = 196.1863  = 49.09164  = 0.3135697  6 / 5
## WARNING: we need the jacobian in the "gradient" attribute
hobbs.ajx <- function(x){
  jj <- hobbs.aj(x)
  jjr <- jj
  attr(jjr, "gradient")<- jj # !!! IMPORTANT
  jjr
}

sagood <- nlfb(x0good, hobbs.res, hobbs.ajx, trace=TRUE)

## no weights
## lower:[1] -Inf -Inf -Inf
## upper:[1] Inf Inf Inf
## Start:lamda: 1e-04  SS= 158.2324  at  = 200  = 50  = 0.3  1 / 0
## <<lamda: 4e-05  SS= 2.61779  at  = 194.3011  = 48.56497  = 0.313994  2 / 1
## <<lamda: 1.6e-05  SS= 2.587325  at  = 196.0825  = 49.07632  = 0.313616  3 / 2
## <<lamda: 6.4e-06  SS= 2.587277  at  = 196.1839  = 49.09134  = 0.313571  4 / 3
## <<lamda: 2.56e-06  SS= 2.587277  at  = 196.1862  = 49.09164  = 0.3135697  5 / 4
## <<lamda: 1.024e-06  SS= 2.587277  at  = 196.1863  = 49.09164  = 0.3135697  6 / 5
```

Test problem – Candlestick

This function was developed by one of us to provide a simple but (for n equal 1 or 2) graphic example of a function with an infinity of solutions for $n \geq 2$. The function can be seen by graphing it to have a spike in the “middle” of a dish, much like some older candlesticks or candle holders. The multiplicity of solutions should make the hessian of a solution singular. For $n = 2$, for example, the minimum lies on a circular locus at the deepest point of the “saucer”.

```
# candlestick function
# J C Nash 2011-2-3
cstick.f<-function(x,alpha=1){
```

```

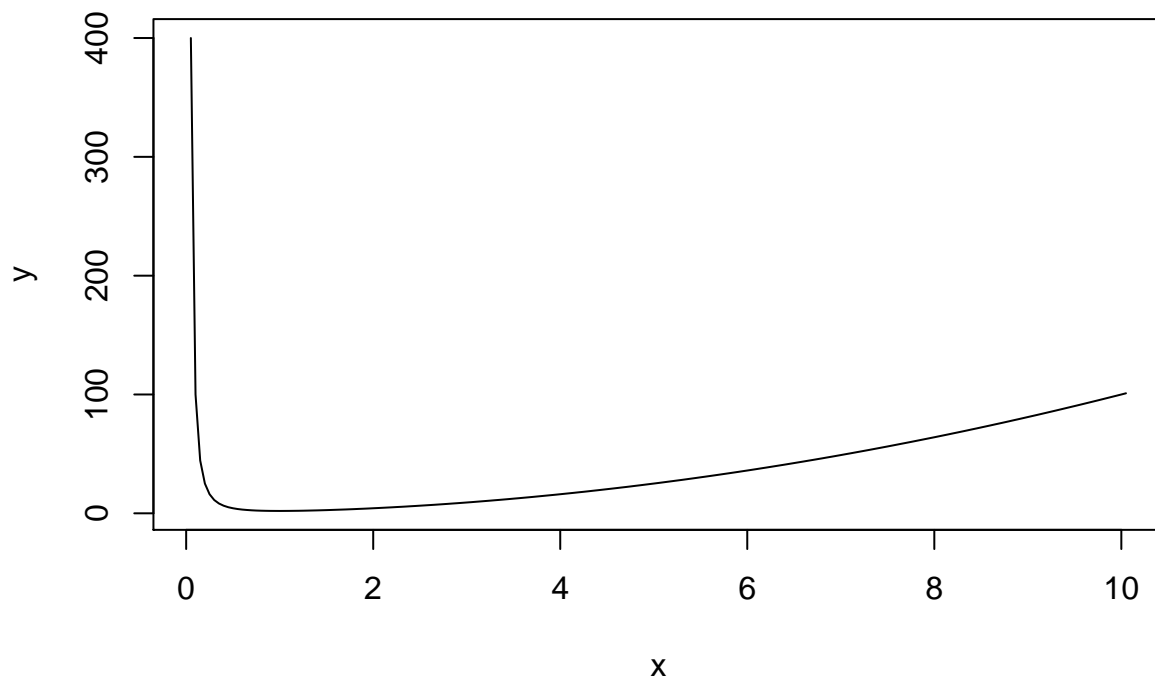
x<-as.vector(x)
r2<-sum(x*x)
f<-as.double(r2+alpha/r2)
return(f)
}

cstick.g<-function(x,alpha=1){
  x<-as.vector(x)
  r2<-sum(x*x)
  g1<-2*x
  g2 <- (-alpha)*2*x/(r2*r2)
  g<-as.double(g1+g2)
  return(g)
}

x <- seq(-100:100)/20.0
y <- x

for (ii in 1:length(x)){
  y[ii] <- cstick.f(x[ii])
}
plot(x, y, type='l') # ?? does not plot from console??

```



```

x0 <- c(1,2)
require(optimx)

## Loading required package: optimx

sdef0 <- optimr(x0, cstick.f, cstick.g, method="Rvmmin", control=list(trace=1))

## Parameter scaling:[1] 1 1
## gradient test tolerance = 6.055454e-06 fval= 5.2
## compare to max(abs(gn-ga))/(1+abs(fval)) = 6.066115e-12

```

```
## Rvminu -- J C Nash 2009-2015 - an R implementation of Alg 21
## Problem of size n= 2   Dot arguments:
## list()
## Initial fn= 5.2
## ig= 1   gnorm= 4.293251      1   1   5.2
## ig= 2   gnorm= 3.884638      2   2   4.468295
## *ig= 3   gnorm= 2.852834      4   3   3.087837
## *ig= 4   gnorm= 2.154552      6   4   2.520416
## ig= 5   gnorm= 3.731383      7   5   2.417429
## ig= 6   gnorm= 1.086155      8   6   2.098346
## ig= 7   gnorm= 0.504457      9   7   2.018112
## ig= 8   gnorm= 0.1371067     10   8   2.001136
## ig= 9   gnorm= 0.01377774    11   9   2.000012
## ig= 10  gnorm= 0.0003470064   12   10  2
## ig= 11  gnorm= 8.985444e-07   13   11  2
## ig= 12  gnorm= 5.846616e-11   14   12  2
## *****No acceptable point
## Converged
## Seem to be done Rvminu
```

```
sdef0
```

```
## $par
## [1] -0.4472136 -0.8944272
##
## $value
## [1] 2
##
## $counts
## function gradient
##      19      12
##
## $convergence
## [1] 0
##
## $message
## [1] "Rvminu appears to have converged"
```

```
xstar <- sdef0$par
gstar <- cstick.g(xstar)
cat("Gradient at proposed solution:")
```

```
## Gradient at proposed solution:
```

```
print(gstar)
```

```
## [1] -2.614686e-11 -5.229372e-11
## FIXED??
## This doesn't seem to work well??
require(autodiffr)
ad_setup()
hc <- autodiffr::hessian(cstick.f)
hstar<-hc(xstar)
cat("Hessian at proposed solution:\n")
```

```
## Hessian at proposed solution:
```

```

print(hstar)

##      [,1] [,2]
## [1,]    0    0
## [2,]    0    0

print(eigen(hstar)$values)

## [1] 0 0
## ?? doesn't seem right
hc(x0)

##      [,1] [,2]
## [1,]    0    0
## [2,]    0    0

require(numDeriv)
hcn0 <- numDeriv::hessian(cstick.f, x0)
hcn0

##      [,1] [,2]
## [1,] 1.984 0.128
## [2,] 0.128 2.176

hcnstar <- numDeriv::hessian(cstick.f, xstar)
hcnstar

##      [,1] [,2]
## [1,]  1.6  3.2
## [2,]  3.2  6.4

hcnj0 <- numDeriv::jacobian(cstick.g, x0)
hcnj0

##      [,1] [,2]
## [1,] 1.984 0.128
## [2,] 0.128 2.176

hcnjstar <- numDeriv::jacobian(cstick.g, xstar)
hcnjstar

##      [,1] [,2]
## [1,]  1.6  3.2
## [2,]  3.2  6.4

eigen(hcnstar)$values

## [1] 8.000000e+00 1.137876e-10

```

Test problem – Wood 4 parameter function

This is reported by Moré, Garbow, and Hillstom (1980) as coming from Colville (1968). The problem in 4 parameters seems to have a false solution far from the accepted one. Is there a good description of this function and the issues it presents?

```

require(autodiffR)
ad_setup() # to ensure it is established
#Example 2: Wood function

```

```

#
wood.f <- function(x){
  res <- 100*(x[1]^2-x[2])^2+(1-x[1])^2+90*(x[3]^2-x[4])^2+(1-x[3])^2+
    10.1*((1-x[2])^2+(1-x[4])^2)+19.8*(1-x[2])*(1-x[4])
  return(res)
}

#gradient:
wood.g <- function(x){
  g1 <- 400*x[1]^3-400*x[1]*x[2]+2*x[1]-2
  g2 <- -200*x[1]^2+220.2*x[2]+19.8*x[4]-40
  g3 <- 360*x[3]^3-360*x[3]*x[4]+2*x[3]-2
  g4 <- -180*x[3]^2+200.2*x[4]+19.8*x[2]-40
  return(c(g1,g2,g3,g4))
}

#hessian:
wood.h <- function(x){
  h11 <- 1200*x[1]^2-400*x[2]+2;    h12 <- -400*x[1]; h13 <- h14 <- 0
  h22 <- 220.2; h23 <- 0;    h24 <- 19.8
  h33 <- 1080*x[3]^2-360*x[4]+2;    h34 <- -360*x[3]
  h44 <- 200.2
  H <- matrix(c(h11,h12,h13,h14,h12,h22,h23,h24,
    h13,h23,h33,h34,h14,h24,h34,h44),ncol=4)
  return(H)
}

#####
x0 <- c(-3,-1,-3,-1) # Wood standard start

cat("Function value at x0=",wood.f(x0),"\n")

## Function value at x0= 19192

wood.ag <- autodiffr::grad(wood.f)
cat("Autodiffr gradient value:")

## Autodiffr gradient value:
vwag0<-wood.ag(x0)
print(vwag0)

## [1] -12008 -2080 -10808 -1880
cat("Manually coded:")

## Manually coded:
vwg0 <- wood.g(x0)
print(vwg0)

## [1] -12008 -2080 -10808 -1880
cat("Differences:\n")

## Differences:
print(vwag0-vwg0)

## [1] 0 0 0 0

```

```

cat("Autodiffr hessian of function value:")

## Autodiffr hessian of function value:
wood.ah <- autodiffr::hessian(wood.f)
vwah0 <- wood.ah(x0)
print(vwah0)

##      [,1] [,2] [,3] [,4]
## [1,] 11202 1200.0    0  0.0
## [2,] 1200  220.2    0 19.8
## [3,]    0   0.0 10082 1080.0
## [4,]    0  19.8  1080  200.2

cat("Autodiffr hessian via jacobian of autodiff gradient value:")

## Autodiffr hessian via jacobian of autodiff gradient value:
wood.ahjag <- autodiffr::jacobian(wood.ag)
vwahjag0<-wood.ahjag(x0)
print(vwahjag0)

##      [,1] [,2] [,3] [,4]
## [1,] 11202 1200.0    0  0.0
## [2,] 1200  220.2    0 19.8
## [3,]    0   0.0 10082 1080.0
## [4,]    0  19.8  1080  200.2

cat("Autodiffr hessian via jacobian of manual gradient value:")

## Autodiffr hessian via jacobian of manual gradient value:
wood.ahj <- autodiffr::jacobian(wood.g)
vwahj0 <- wood.ah(x0)
print(vwahj0)

##      [,1] [,2] [,3] [,4]
## [1,] 11202 1200.0    0  0.0
## [2,] 1200  220.2    0 19.8
## [3,]    0   0.0 10082 1080.0
## [4,]    0  19.8  1080  200.2

cat("Manually coded:")

## Manually coded:
vwah0<-wood.h(x0)
print(vwah0)

##      [,1] [,2] [,3] [,4]
## [1,] 11202 1200.0    0  0.0
## [2,] 1200  220.2    0 19.8
## [3,]    0   0.0 10082 1080.0
## [4,]    0  19.8  1080  200.2

cat("Differences from vwah0\n")

## Differences from vwah0

```



```

cat("vwah0\n")

## vwah0
print(vwah0-vwh0)

##      [,1] [,2] [,3] [,4]
## [1,]    0    0    0    0
## [2,]    0    0    0    0
## [3,]    0    0    0    0
## [4,]    0    0    0    0

cat("\n")

cat("vwahj0\n")

## vwahj0
print(vwahj0-vwh0)

##      [,1] [,2] [,3] [,4]
## [1,]    0    0    0    0
## [2,]    0    0    0    0
## [3,]    0    0    0    0
## [4,]    0    0    0    0

cat("\n")

cat("vwahjag0\n")

## vwahjag0
print(vwahjag0-vwh0)

##      [,1] [,2] [,3] [,4]
## [1,]    0    0    0    0
## [2,]    0    0    0    0
## [3,]    0    0    0    0
## [4,]    0    0    0    0

cat("\n")

## d <- c(1,1,1,1)
require(optimx)
meths <- c("snewton", "snewtonm", "nlm")
wdefault <- opm(x0, fn=wood.f, gr=wood.g, hess=wood.h, method=meths, control=list(trace=0))
print(wdefault)

##      p1 p2 p3 p4      value fevals gevals convergence kkt1 kkt2
## snewton   1  1  1  1 1.142616e-29   119    70          92 TRUE TRUE
## snewtonm  1  1  1  1 1.399599e-28    88    50           0 TRUE TRUE
## nlm       1  1  1  1 1.004943e-16    NA   335           0 TRUE TRUE
##      xtime
## snewton  0.008
## snewtonm 0.004
## nlm      0.008

wagah <- opm(x0, fn=wood.f, gr=wood.ag, hess=wood.ah, method=meths, control=list(trace=0))

## Small gradient

```

```
print(wagah)
```

```
##          p1 p2 p3 p4          value fevals gevals convergence kkt1 kkt2
## snewton   1  1  1  1  0.000000e+00    116    67           0 TRUE TRUE
## snewtonm  1  1  1  1  0.000000e+00    81    50           0 TRUE TRUE
## nlm       1  1  1  1  1.004942e-16     NA   335           0 TRUE TRUE
##          xtime
## snewton   4.388
## snewtonm  3.324
## nlm       24.740
```

Performance issues

Optimization is, by its very nature, about improving things. Thus it is of prime interest to seek faster and better ways to optimize functions. In this section we look at some issues that may influence the speed, reliability and correctness of optimization calculations.

First, it is critical to note that **R** almost always offers several ways to accomplish the same computational result. However, the speed with which the different approaches return a result can be wildly different. (?? can JN find the 800% scale factor example??).

Second, there are many parts of the autodiff wrapper of Julia's automatic differentiation that may use up computing cycles:

- We must translate from one programming language to another in some sense in order to call the appropriate functions in Julia based on **R** functions.
- Results must be properly structured on return to **R**.
- Hand coded derivative expressions, especially hand-optimized ones, can be expected to out-perform automatic differentiation results.

NOTE: Performance is interesting, but it is far from the complete picture. We can use results from autodiff to validate hand-coded functions. We can get results that are efficient of human time and effort that may be otherwise unavailable. Moreover, the results of computing gradients and Hessians allow us to conclude that a solution has been achieved.

A small performance comparison using autodiff

```
rm(list=ls())
require(autodiff)
autodiff::ad_setup() # to ensure it is established

ores <- function(x){
  x # Function will be the parameters. ofn is sum of squares
}

logit <- function(x) exp(x) / (1 + exp(x))

ofn <- function(x){
  res <- ores(x) # returns a vector of residual values
  sum(logit(res) ^ 2)
}
```

```

## Now try to generate the gradient function
ogr <- autodiffr::grad(ofn)

system.time(ogr(runif(100)))

##      user  system elapsed
##    0.232   0.000   0.235

system.time(ogr(runif(100)))

##      user  system elapsed
##    0.02    0.00    0.02

ogr1 <- autodiffr::grad(ofn, xsize = runif(100))

system.time(ogr1(runif(100)))

##      user  system elapsed
##    0.012   0.000   0.013

system.time(ogr1(runif(100)))

##      user  system elapsed
##    0.012   0.000   0.013

ogr2 <- autodiffr::grad(ofn, xsize = runif(100), use_tape = TRUE)

system.time(ogr2(runif(100)))

##      user  system elapsed
##    0.136   0.000   0.135

system.time(ogr2(runif(100)))

##      user  system elapsed
##    0.008   0.000   0.006

```

Bibliography

- Colville, A R. 1968. “A Comparative Study of Nonlinear Programming Codes, Rep. 320-2949.” IBM New York Scientific Center.
- Fletcher, R. 1965. “Function Minimization Without Calculating Derivatives – a Review.” *Computer Journal* 8: 33–41.
- Moré, J. J., B. S. Garbow, and K. E. Hillstom. 1980. “ANL-80-74, User Guide for MINPACK-1.” Argonne National Laboratory. <http://www.mcs.anl.gov/~jmore/ANL8074a.pdf>.
- Nash, John C. 1979. *Compact Numerical Methods for Computers: Linear Algebra and Function Minimisation*. Bristol: Adam Hilger.