

# Optimization and related uses of autodiff: Illustrations

*Changcheng Li and John C. Nash*

*2018/7/10*

## Introduction

**autodiff** is an **R** package to perform automatic differentiation of **R** functions by calling the automatic differentiation tools in the **Julia** programming language. The document **FandR** (??ref) describes how to install **autodiff**.

Here we will illustrate how **autodiff** can be used to provide gradients or other derivative information for use in optimization problems for which **R** is being used to attempt solutions.

## Problem setup

Most methods for optimization require

- an objective function that is to be minimized or maximized. Because the minimum of  $f(x)$  is the maximum of  $-f(x)$ , we will only talk of minimizing functions. Though the mathematics of this are trivial, the care and attention to avoid errors when translating a maximization problem to one involving minimization require serious effort and continuous checking.
- a starting set of values for the parameters to be optimized (and perhaps also for any exogenous data and/or fixed parameters)

We need to load the **autodiff** package and then initiate it. NOTE: This is quite slow the first time it is run. WORSE: It is always treated as a “first time” when called in knitr during the processing of a vignette from an Rmd-type file.

```
library(autodiff)
```

```
##
## Attaching package: 'autodiff'
## The following object is masked from 'package:stats':
##
##      deriv
ad_setup()
```

```
## Julia version 0.6.3 at location /usr/local/bin will be used.
## Loading setup script for JuliaCall...
## Finish loading setup script for JuliaCall.
```

And we can use package **numDeriv** to compare with **autodiff**.

```
require(numDeriv)
```

```
## Loading required package: numDeriv
##
## Attaching package: 'numDeriv'
## The following objects are masked from 'package:autodiff':
##
##      grad, hessian, jacobian
```

## Test problem – ViaRes

This is simply to test how we can get the gradient of a function that is defined as the sum of squares of residuals, BUT the residuals are computed in a subsidiary function that must be called.

At July 2, 2018, this gives an error that stops knitr, so evaluation is turned off in the following example.

```
require(autodiffr)
ad_setup() # to ensure it is established

ores <- function(x){
  x # Function will be the parameters. ofn is sum of squares
}

ofn0 <- function(x){ # original ofn
  res <- ores(x) # returns a vector of residual values
  val <- as.numeric(crossprod(res)) # as.numeric because crossprod is a matrix
  val
}

ofn <- function(x){ # But autodiffr does not understand crossprod()
  res <- ores(x) # returns a vector of residual values
  val <- sum(res*res) # NOT crossprod()
  val
}
```

Note that this works with eval=TRUE, but Chebyquad still failing.

```
## Now try to generate the gradient function
ogr <- autodiffr::makeGradFunc(ofn)

# print(ogr) # this will be more or less meaningless link to Julia function
x0 <- c(1,2,3)
print(ofn(x0)) # should be 14
```

```
## [1] 14
```

```
print(ofn0(x0)) # should be 14
```

```
## [1] 14
```

```
ogr0<-ogr(x0) # should be 2, 4, 6
ogr0
```

```
## [1] 2 4 6
```

## Test problem – Chebyquad

This problem was given prominence in the optimization literature by Fletcher (1965).

First let us define our Chebyquad function. Note that this is for the **vector** x. This version is expressed as the sum of squares of a vector of function values, which provide a nonlinear least squares problem. Note that `crossprod()` may cause difficulties as it is not written in **R**.

```
require(autodiffr)
ad_setup()

cyq.f <- function (x) {
```

```

rv<-cyq.res(x)
f <- sum(rv*rv)
}

cyq.res <- function (x) {
# Fletcher's chebyquad function m = n -- residuals
  n<-length(x)
  res<-zeros(x) # need to use zeros() to do initialization instead of rep() otherwise autodiffr will b
  ## This is because later on res[i] <- rr,
  ## if res is a normal R vector and rr is some Julia Number used by autodiffr,
  ## then R doesn't know what to do.
  for (i in 1:n) { #loop over resid
    rr<-0.0
    for (k in 1:n) {
      z7<-1.0
      z2<-2.0*x[k]-1.0
      z8<-z2
      j<-1
      while (j<i) {
        z6<-z7
        z7<-z8
        z8<-2*z2*z7-z6 # recurrence to compute Chebyshev polynomial
        j<-j+1
      } # end recurrence loop
      rr<-rr+z8
    } # end loop on k
    rr<-rr/n
    if (2*trunc(i/2) == i) { rr <- rr + 1.0/(i*i - 1) }
    res[i]<-rr
  } # end loop on i
  res
}

```

Let us choose a single value for the number of parameters, and for illustration use  $n = 4$ .

```

## cyq.setup
n <- 4
lower<-rep(-10.0, n)
upper<-rep(10.0, n)
x<-1:n
x<-x/(n+1.0) # Initial value suggested by Fletcher

```

For safety, let us check the function and a numerical approximation to the gradient.

```

require(numDeriv)
cat("Initial parameters:")

## Initial parameters:
print(x)

## [1] 0.2 0.4 0.6 0.8
cat("Initial value of the function is ",cyq.f(x),"\n")

## Initial value of the function is 0.07118393

```

```
gn <- numDeriv::grad(cyq.f, x) # using numDeriv
cat("Approximation to gradient at initial point:")
```

```
## Approximation to gradient at initial point:
```

```
print(gn)
```

```
## [1] 0.6170624 0.1882112 -0.1882112 -0.6170624
```

Using a modular approach to the problem, first specifying it via **residuals** and computing the function as a sum of squares, we can also generate the gradient.

```
# Ref: Fletcher, R. (1965) Function minimization without calculating derivatives -- a review,
#       Computer J., 8, 33-41.
```

```
# Note we do not have all components here e.g., .jsd, .h
```

```
cyq.jac<- function (x) {
# Chebyquad Jacobian matrix
n<-length(x)
cj<-matrix(0.0, n, n)
for (i in 1:n) { # loop over rows
  for (k in 1:n) { # loop over columns (parameters)
    z5<-0.0
    cj[i,k]<-2.0
    z8<-2.0*x[k]-1.0
    z2<-z8
    z7<-1.0
    j<- 1
    while (j<i) { # recurrence loop
      z4<-z5
      z5<-cj[i,k]
      cj[i,k]<-4.0*z8+2.0*z2*z5-z4
      z6<-z7
      z7<-z8
      z8<-2.0*z2*z7-z6
      j<- j+1
    } # end recurrence loop
    cj[i,k]<-cj[i,k]/n
  } # end loop on k
} # end loop on i
cj
}
```

```
cyq.g <- function (x) {
  cj<-cyq.jac(x)
  rv<-cyq.res(x)
  gg<- 2.0 * as.vector(rv %*% cj)
}
```

```
# check gradient function cyq.g
```

```
gajn <- cyq.g(x)
```

```
print(gajn)
```

```
## [1] 0.6170624 0.1882112 -0.1882112 -0.6170624
```

We can now try to see if `autodiffr` matches this gradient. However, the following code gives an error in Julia.

```
# Do not evaluate, as this fails # Now it should work
cyq.ag <- autodiffr::makeGradFunc(cyq.f)
gaag <- cyq.ag(x)
print(gaag)
```

```
## [1] 0.6170624 0.1882112 -0.1882112 -0.6170624
```

As a workaround, we can get the Chebyquad function from the package `funconstrain`. The `funconstrain` offering does NOT require a call to the residuals, but has a single level **R** function.

```
require(funconstrain)
```

```
## Loading required package: funconstrain
```

```
cat("funconstrain loaded\n")
```

```
## funconstrain loaded
```

```
cheb <- chebyquad() # Seem to need the brackets or doesn't return pieces
print(str(cheb))
```

```
## List of 4
```

```
## $ fn:function (par)
```

```
## $ gr:function (par)
```

```
## $ fg:function (par)
```

```
## $ x0:function (n = 50)
```

```
## NULL
```

```
cyq2.f <- cheb$fn
```

```
## Note that funconstrain offers the starting value
```

```
## x0b <- cheb$x0(n=4) # Need the size of the vector
```

```
## x0b
```

```
## cyq2.f(x0b)
```

```
## same as
```

```
print(cyq2.f(x))
```

```
## [1] 0.07118393
```

```
## Try the gradient
```

```
cyq2.ag <- autodiffr::makeGradFunc(cyq2.f) # Need autodiffr:: specified for knitr
```

```
## print(cyq2.g)
```

```
cat("Gradient at x")
```

```
## Gradient at x
```

```
g2ag <- cyq2.ag(x)
```

```
print(g2ag)
```

```
## [1] 0.6170624 0.1882112 -0.1882112 -0.6170624
```

```
require(microbenchmark)
```

```
## Loading required package: microbenchmark
```

```
# NOTE: Very strange output when running 1 line at a time here in Rstudio
cat("cyq.f timing:\n")
```

```
## cyq.f timing:
```

```
tcyq.f <- microbenchmark(cyq.f(x))
tcyq.f
```

```
## Unit: microseconds
```

```
##      expr      min       lq      mean  median       uq      max neval
##  cyq.f(x) 386.105 421.814 435.8203 427.4625 435.1135 1132.156   100
```

```
cat("cyq2.f timing:\n")
```

```
## cyq2.f timing:
```

```
tcyq2.f <- microbenchmark(cyq2.f(x))
summary(tcyq2.f)
```

```
##      expr      min       lq      mean median       uq      max neval
## 1 cyq2.f(x)  6.232  6.4605  6.99143   6.632  6.8575  33.252   100
```

```
tcyq.g <- microbenchmark(cyq.g(x))
summary(tcyq.g)
```

```
##      expr      min       lq      mean  median       uq      max neval
## 1 cyq.g(x) 421.028 432.098 508.6729 461.5245 477.2615 5580.831   100
```

```
tcyq.g <- microbenchmark(cyq.g(x))
tcyq.g
```

```
## Unit: microseconds
```

```
##      expr      min       lq      mean  median       uq      max neval
##  cyq.g(x) 421.218 436.4315 465.3071 455.4265 473.7935 1101.629   100
```

```
cyq2.g <- cheb$gr
tcyq2.g <- microbenchmark(cyq2.g(x), unit="us" )
# microseconds
tcyq2.g
```

```
## Unit: microseconds
```

```
##      expr      min       lq      mean  median       uq      max neval
##  cyq2.g(x) 12.548 12.9735 14.03828 13.3575 14.0655  39.024   100
```

```
# These are very slow
tcyq.ag <- microbenchmark(cyq.ag(x), unit="us" )
# microseconds
tcyq.ag
```

```
## Unit: microseconds
```

```
##      expr      min       lq      mean  median       uq      max neval
##  cyq.ag(x) 99669.59 101510.8 109016.6 105695.8 113695.7 148397.2   100
```

```
tcyq2.ag <- microbenchmark(cyq2.ag(x), unit="us" )
# microseconds
tcyq2.ag
```

```
## Unit: microseconds
```

```
##      expr      min       lq      mean  median       uq      max neval
##  cyq2.ag(x) 21258.3 21617.71 22583.94 21999.72 22408.28 39148.11   100
```

```

# These are quicker, but still slow
cyq.optimized_ag <- autodiffr::makeGradFunc(cyq.f, x = runif(length(x)), use_tape = TRUE)

cyq2.optimized_ag <- autodiffr::makeGradFunc(cyq2.f, x = runif(length(x)), use_tape = TRUE)

tcyq.optimized_ag <- microbenchmark(cyq.optimized_ag(x), unit="us" )
# microseconds
tcyq.optimized_ag

```

```

## Unit: microseconds
##          expr      min       lq      mean    median       uq      max
##  cyq.optimized_ag(x) 897.415 903.732 4714.804 910.4105 944.7425 378929.7
##    neval
##      100

```

```

tcyq2.optimized_ag <- microbenchmark(cyq2.optimized_ag(x), unit="us" )
# microseconds
tcyq2.optimized_ag

```

```

## Unit: microseconds
##          expr      min       lq      mean    median       uq      max
##  cyq2.optimized_ag(x) 868.883 883.8445 2861.381 913.7315 940.917 194449.4
##    neval
##      100

```

```

## The slowness of the optimized method is partly because the overhead of the JuliaCall and autodiffr p
## After interface functions become stable, I will try to carry on some performance optimizations,
## which is a goal of the project at last phase.
## For example, even we only have a very simple function, the timing is high because of the overhead.
foobar <- function(x) sum(x)

```

```

foobar.ag <- autodiffr::makeGradFunc(foobar, x = runif(length(x)), use_tape = TRUE)

tfoobar.ag <- microbenchmark(foobar.ag(x), unit="us" )
tfoobar.ag

```

```

## Unit: microseconds
##          expr      min       lq      mean    median       uq      max neval
##   foobar.ag(x) 843.367 852.1865 1037.841 859.5525 878.765 16446.05   100

```

```

## Suppose we are dealing with input of larger size, the overhead stays roughly the same,
## so the overhead should matters not that much as in the case n = 4.
## For example, if n = 25, the diffrence of performance in ratio is not that much.
## cyq.setup
n <- 25
lower<-rep(-10.0, n)
upper<-rep(10.0, n)
x<-1:n
x<-x/(n+1.0) # Initial value suggested by Fletcher

```

```

tcyq.g <- microbenchmark(cyq.g(x), unit = "us")
tcyq.g

```

```

## Unit: microseconds
##          expr      min       lq      mean    median       uq      max neval
##   cyq.g(x) 6407.557 6506.096 6813.755 6630.014 6858.59 12014.49   100

```

```

tcyq2.g <- microbenchmark(cyq2.g(x), unit="us" )
# microseconds
tcyq2.g

## Unit: microseconds
##      expr      min       lq      mean  median       uq      max neval
##  cyq2.g(x) 72.005 74.5475 77.59262 75.763 78.9035 107.534   100
## The bad thing is that if the input size changes, we need to make an optimized gradient again.

cyq.optimized_ag <- autodiffr::makeGradFunc(cyq.f, x = runif(length(x)), use_tape = TRUE)
cyq2.optimized_ag <- autodiffr::makeGradFunc(cyq2.f, x = runif(length(x)), use_tape = TRUE)

tcyq.optimized_ag <- microbenchmark(cyq.optimized_ag(x), unit="us" )
# microseconds
tcyq.optimized_ag

## Unit: microseconds
##      expr      min       lq      mean  median       uq      max
##  cyq.optimized_ag(x) 11472.5 12016.28 12643.18 12306.49 12754.31 16387.49
##    neval
##      100

tcyq2.optimized_ag <- microbenchmark(cyq2.optimized_ag(x), unit="us" )
# microseconds
tcyq2.optimized_ag

## Unit: microseconds
##      expr      min       lq      mean  median       uq      max
##  cyq2.optimized_ag(x) 937.189 958.2585 1249.537 1002.124 1058.863 23098.52
##    neval
##      100

## Also note it is better to check the correctness when generating optimized gradient,
all.equal(cyq.g(x), cyq.optimized_ag(x))

## [1] TRUE
all.equal(cyq2.g(x), cyq2.optimized_ag(x))

## [1] TRUE
## Benchmarking times without user interface wrappers
tape1 <- reverse.grad.tape(cyq.f, runif(length(x)))
microbenchmark(reverse.grad(tape1, x), unit="us" )

## Unit: microseconds
##      expr      min       lq      mean  median       uq
##  reverse.grad(tape1, x) 11650.72 12386.13 12887.6 12694.86 13185.99
##      max neval
## 17501.4   100

tape2 <- reverse.grad.tape(cyq2.f, runif(length(x)))
microbenchmark(reverse.grad(tape2, x), unit="us" )

## Unit: microseconds
##      expr      min       lq      mean  median       uq      max
##  reverse.grad(tape2, x) 929.1 934.066 965.2175 939.575 951.16 1724.895

```



```

## neval
## 100

## Benchmarking times without autodiffr wrappers
JuliaCall::julia_command("using ReverseDiff")
tape1 <- reverse.grad.tape(cyq.f, runif(length(x)))
microbenchmark(JuliaCall::julia_call("ReverseDiff.gradient!", tape1, x), unit="us" )

## Unit: microseconds
##                                     expr      min
## JuliaCall::julia_call("ReverseDiff.gradient!", tape1, x) 12135.08
##      lq      mean  median      uq      max neval
## 13220.72 13715.31 13587.86 14208.88 16231.08   100

tape2 <- reverse.grad.tape(cyq2.f, runif(length(x)))
microbenchmark(JuliaCall::julia_call("ReverseDiff.gradient!", tape2, x), unit="us" )

## Unit: microseconds
##                                     expr      min      lq
## JuliaCall::julia_call("ReverseDiff.gradient!", tape2, x) 555.604 560.7235
##      mean  median      uq      max neval
## 593.0419 564.0385 594.4725 1392.169   100

```

## Test problem – Hobbs weed infestation

This nonlinear estimation problem was brought to one of the authors (JN) in the mid-1970s (See Nash (1979)). It has just 12 data points and asks for the estimation of a 3-parameter logistic growth curve. The present example does not provide for scaling.

```

hobbs.f<- function(x){ # Hobbs weeds problem -- function
  if (abs(12*x[3]) > 500) { # check computability
    fbad<-.Machine$double.xmax
    return(fbad)
  }
  res<-hobbs.res(x)
  f<-sum(res*res)
}

hobbs.res<-function(x){ # Hobbs weeds problem -- residual
# This variant uses looping
  if(length(x) != 3) stop("hobbs.res -- parameter vector n!=3")
  y<-c(5.308, 7.24, 9.638, 12.866, 17.069, 23.192, 31.443, 38.558, 50.156, 62.948,
      75.995, 91.972)
  t<-1:12
  if(abs(12*x[3])>50) {
    res<-rep(Inf,12)
  } else {
    res<-x[1]/(1+x[2]*exp(-x[3]*t)) - y
  }
}

hobbs.jac<-function(x){ # Jacobian of Hobbs weeds problem
  jj<-matrix(0.0, 12, 3)
  t<-1:12

```

```

yy<-exp(-x[3]*t)
zz<-1.0/(1+x[2]*yy)
jj[t,1] <- zz
jj[t,2] <- -x[1]*zz*zz*yy
jj[t,3] <- x[1]*zz*zz*yy*x[2]*t
jjret <- jj
attr(jjret,"gradient") <- jj
return(jjret)
}

hobbs.g<-function(x){ # gradient of Hobbs weeds problem
# NOT EFFICIENT TO CALL AGAIN
jj<-hobbs.jac(x)
res<-hobbs.res(x)
gg<-as.vector(2.*t(jj) %*% res)
return(gg)
}

hobbs.rsd<-function(x) { # Jacobian second derivative
rsd<-array(0.0, c(12,3,3))
t<-1:12
yy<-exp(-x[3]*t)
zz<-1.0/(1+x[2]*yy)
rsd[t,1,1]<- 0.0
rsd[t,2,1]<- -yy*zz*zz
rsd[t,1,2]<- -yy*zz*zz
rsd[t,2,2]<- 2.0*x[1]*yy*yy*zz*zz*zz
rsd[t,3,1]<- t*x[2]*yy*zz*zz
rsd[t,1,3]<- t*x[2]*yy*zz*zz
rsd[t,3,2]<- t*x[1]*yy*zz*zz*(1-2*x[2]*yy*zz)
rsd[t,2,3]<- t*x[1]*yy*zz*zz*(1-2*x[2]*yy*zz)
## rsd[t,3,3]<- 2*t*t*x[1]*x[2]*x[2]*yy*yy*zz*zz*zz
rsd[t,3,3]<- -t*t*x[1]*x[2]*yy*zz*zz*(1-2*yy*zz*x[2])
return(rsd)
}

hobbs.h <- function(x) { ## compute Hessian
# cat("Hessian not yet available\n")
# return(NULL)
H<-matrix(0,3,3)
res<-hobbs.res(x)
jj<-hobbs.jac(x)
rsd<-hobbs.rsd(x)
## H<-2.0*(t(res) %*% rsd + t(jj) %*% jj)
for (j in 1:3) {
for (k in 1:3) {
for (i in 1:12) {
H[j,k]<-H[j,k]+res[i]*rsd[i,j,k]
}
}
}
}
H<-2*(H + t(jj) %*% jj)

```

```

    return(H)
}

x0good <- c(200, 50, 0.3)
x0bad <- c(1,1,1)
f0good <- hobbs.f(x0good)
cat("Sum of squares at the GOOD starting point:",f0good,"\n")

## Sum of squares at the GOOD starting point: 158.2324
f0bad <- hobbs.f(x0bad)
cat("Sum of squares at the BAD starting point:",f0bad,"\n")

## Sum of squares at the BAD starting point: 23520.58
res0good <- hobbs.res(x0good)
## Residuals -- good starting point
res0good

## [1] -0.05050236 -0.20779506 -0.26086802 -0.41247546 -0.61690702
## [6] -1.60525418 -3.36423901 -2.43016453 -4.28734016 -5.63079076
## [11] -5.67542775 -7.44779537
res0bad <- hobbs.res(x0bad)
## Residuals -- bad starting point
res0bad

## [1] -4.576941 -6.359203 -8.685426 -11.883986 -16.075693 -22.194473
## [7] -30.443911 -37.558335 -49.156123 -61.948045 -74.995017 -90.972006

require(autodiffr)
ad_setup()
hobbs.ag <- autodiffr::makeGradFunc(hobbs.f)
hobbsag0good <- hobbs.ag(x0good)
## Gradient by AD -- good starting point
hobbsag0good

## [1] -17.8438 48.2491 -24559.8187
## Compare hand coded function
hobbsggood <- hobbs.g(x0good)
hobbsggood

## [1] -17.8438 48.2491 -24559.8187
## Gradient by AD -- bad starting point
hobbsag0bad <- hobbs.ag(x0bad)
hobbsag0bad

## [1] -824.042084 4.764888 -11.025384
## Compare hand coded function
hobbsgbad <- hobbs.g(x0bad)
hobbsgbad

## [1] -824.042084 4.764888 -11.025384
## Interestingly, the magnitude of gradient elements greater for "good"

```

```
hobbs.aj <- autodiffr::makeJacobianFunc(hobbs.res)
## Gradient by AD -- good starting point
hobbsaj0good <- hobbs.aj(x0good)
hobbsaj0good
```

```
##           [,1]      [,2]      [,3]
## [1,] 0.02628749 -0.1023858  5.119291
## [2,] 0.03516102 -0.1356989 13.569891
## [3,] 0.04688566 -0.1787496 26.812437
## [4,] 0.06226762 -0.2335615 46.712293
## [5,] 0.08226046 -0.3019747 75.493681
## [6,] 0.10793373 -0.3851362 115.540847
## [7,] 0.14039380 -0.4827335 168.956738
## [8,] 0.18063918 -0.5920347 236.813864
## [9,] 0.22934330 -0.7069798 318.140911
## [10,] 0.28658605 -0.8178179 408.908969
## [11,] 0.35159786 -0.9119072 501.548971
## [12,] 0.42262102 -0.9760500 585.629985
```

```
## Compare hand coded function
hobbsjgood <- hobbs.jac(x0good)
hobbsjgood
```

```
##           [,1]      [,2]      [,3]
## [1,] 0.02628749 -0.1023858  5.119291
## [2,] 0.03516102 -0.1356989 13.569891
## [3,] 0.04688566 -0.1787496 26.812437
## [4,] 0.06226762 -0.2335615 46.712293
## [5,] 0.08226046 -0.3019747 75.493681
## [6,] 0.10793373 -0.3851362 115.540847
## [7,] 0.14039380 -0.4827335 168.956738
## [8,] 0.18063918 -0.5920347 236.813864
## [9,] 0.22934330 -0.7069798 318.140911
## [10,] 0.28658605 -0.8178179 408.908969
## [11,] 0.35159786 -0.9119072 501.548971
## [12,] 0.42262102 -0.9760500 585.629985
## attr("gradient")
##           [,1]      [,2]      [,3]
## [1,] 0.02628749 -0.1023858  5.119291
## [2,] 0.03516102 -0.1356989 13.569891
## [3,] 0.04688566 -0.1787496 26.812437
## [4,] 0.06226762 -0.2335615 46.712293
## [5,] 0.08226046 -0.3019747 75.493681
## [6,] 0.10793373 -0.3851362 115.540847
## [7,] 0.14039380 -0.4827335 168.956738
## [8,] 0.18063918 -0.5920347 236.813864
## [9,] 0.22934330 -0.7069798 318.140911
## [10,] 0.28658605 -0.8178179 408.908969
## [11,] 0.35159786 -0.9119072 501.548971
## [12,] 0.42262102 -0.9760500 585.629985
```

```
## Gradient by AD -- bad starting point
hobbsaj0bad <- hobbs.aj(x0bad)
hobbsaj0bad
```

```
##           [,1]           [,2]           [,3]
## [1,] 0.7310586 -1.966119e-01 1.966119e-01
## [2,] 0.8807971 -1.049936e-01 2.099872e-01
## [3,] 0.9525741 -4.517666e-02 1.355300e-01
## [4,] 0.9820138 -1.766271e-02 7.065082e-02
## [5,] 0.9933071 -6.648057e-03 3.324028e-02
## [6,] 0.9975274 -2.466509e-03 1.479906e-02
## [7,] 0.9990889 -9.102212e-04 6.371548e-03
## [8,] 0.9996646 -3.352377e-04 2.681901e-03
## [9,] 0.9998766 -1.233793e-04 1.110414e-03
## [10,] 0.9999546 -4.539581e-05 4.539581e-04
## [11,] 0.9999833 -1.670114e-05 1.837126e-04
## [12,] 0.9999939 -6.144137e-06 7.372964e-05
```

```
## Compare hand coded function
```

```
hobbsjbad <- hobbs.jac(x0bad)
hobbsjbad
```

```
##           [,1]           [,2]           [,3]
## [1,] 0.7310586 -1.966119e-01 1.966119e-01
## [2,] 0.8807971 -1.049936e-01 2.099872e-01
## [3,] 0.9525741 -4.517666e-02 1.355300e-01
## [4,] 0.9820138 -1.766271e-02 7.065082e-02
## [5,] 0.9933071 -6.648057e-03 3.324028e-02
## [6,] 0.9975274 -2.466509e-03 1.479906e-02
## [7,] 0.9990889 -9.102212e-04 6.371548e-03
## [8,] 0.9996646 -3.352377e-04 2.681901e-03
## [9,] 0.9998766 -1.233793e-04 1.110414e-03
## [10,] 0.9999546 -4.539581e-05 4.539581e-04
## [11,] 0.9999833 -1.670114e-05 1.837126e-04
## [12,] 0.9999939 -6.144137e-06 7.372964e-05
## attr("gradient")
##           [,1]           [,2]           [,3]
## [1,] 0.7310586 -1.966119e-01 1.966119e-01
## [2,] 0.8807971 -1.049936e-01 2.099872e-01
## [3,] 0.9525741 -4.517666e-02 1.355300e-01
## [4,] 0.9820138 -1.766271e-02 7.065082e-02
## [5,] 0.9933071 -6.648057e-03 3.324028e-02
## [6,] 0.9975274 -2.466509e-03 1.479906e-02
## [7,] 0.9990889 -9.102212e-04 6.371548e-03
## [8,] 0.9996646 -3.352377e-04 2.681901e-03
## [9,] 0.9998766 -1.233793e-04 1.110414e-03
## [10,] 0.9999546 -4.539581e-05 4.539581e-04
## [11,] 0.9999833 -1.670114e-05 1.837126e-04
## [12,] 0.9999939 -6.144137e-06 7.372964e-05
```

Now let us try this in a solution of nonlinear least squares.

WARNING: Because of some compatibility issues with other **R** software, the jacobian must be available in the “gradient” attribute returned by the jacobian function. The purpose of this is to allow the function `nlshr::nlfb` to have the same name for the residual and jacobian function. This is used in generating a symbolic jacobian function in `nlshr::nlxb`. However, it can catch unwary users (including us!).

```
# try in a function
```

```
require(nlshr)
```

```
## Loading required package: nlshr
```

```

## manual
smgood <- nlfb(x0good, hobbs.res, hobbs.jac, trace=TRUE)

## no weights
## lower:[1] -Inf -Inf -Inf
## upper:[1] Inf Inf Inf
## Start:lamda: 1e-04 SS= 158.2324 at = 200 = 50 = 0.3 1 / 0
## <<lamda: 4e-05 SS= 2.61779 at = 194.3011 = 48.56497 = 0.313994 2 / 1
## <<lamda: 1.6e-05 SS= 2.587325 at = 196.0825 = 49.07632 = 0.313616 3 / 2
## <<lamda: 6.4e-06 SS= 2.587277 at = 196.1839 = 49.09134 = 0.313571 4 / 3
## <<lamda: 2.56e-06 SS= 2.587277 at = 196.1862 = 49.09164 = 0.3135697 5 / 4
## <<lamda: 1.024e-06 SS= 2.587277 at = 196.1863 = 49.09164 = 0.3135697 6 / 5

## WARNING: we need the jacobian in the "gradient" attribute
hobbs.ajx <- function(x){
  jj <- hobbs.aj(x)
  jjr <- jj
  attr(jjr, "gradient")<- jjr ### IMPORTANT
  jjr
}

sagood <- nlfb(x0good, hobbs.res, hobbs.ajx, trace=TRUE)

## no weights
## lower:[1] -Inf -Inf -Inf
## upper:[1] Inf Inf Inf
## Start:lamda: 1e-04 SS= 158.2324 at = 200 = 50 = 0.3 1 / 0
## <<lamda: 4e-05 SS= 2.61779 at = 194.3011 = 48.56497 = 0.313994 2 / 1
## <<lamda: 1.6e-05 SS= 2.587325 at = 196.0825 = 49.07632 = 0.313616 3 / 2
## <<lamda: 6.4e-06 SS= 2.587277 at = 196.1839 = 49.09134 = 0.313571 4 / 3
## <<lamda: 2.56e-06 SS= 2.587277 at = 196.1862 = 49.09164 = 0.3135697 5 / 4
## <<lamda: 1.024e-06 SS= 2.587277 at = 196.1863 = 49.09164 = 0.3135697 6 / 5

```

## Test problem – Candlestick

This function was developed by one of us to provide a simple but (for  $n$  equal 1 or 2) graphic example of a function with an infinity of solutions for  $n \geq 2$ . The function can be seen by graphing it to have a spike in the “middle” of a dish, much like some older candlesticks or candle holders. The multiplicity of solutions should make the hessian of a solution singular. For  $n = 2$ , for example, the minimum lies on a circular locus at the deepest point of the “saucer”.

```

# candlestick function
# J C Nash 2011-2-3
cstick.f<-function(x,alpha=1){
  x<-as.vector(x)
  r2<-sum(x*x)
  f<-as.double(r2+alpha/r2)
  return(f)
}

cstick.g<-function(x,alpha=1){
  x<-as.vector(x)
  r2<-sum(x*x)
  g1<-2*x

```

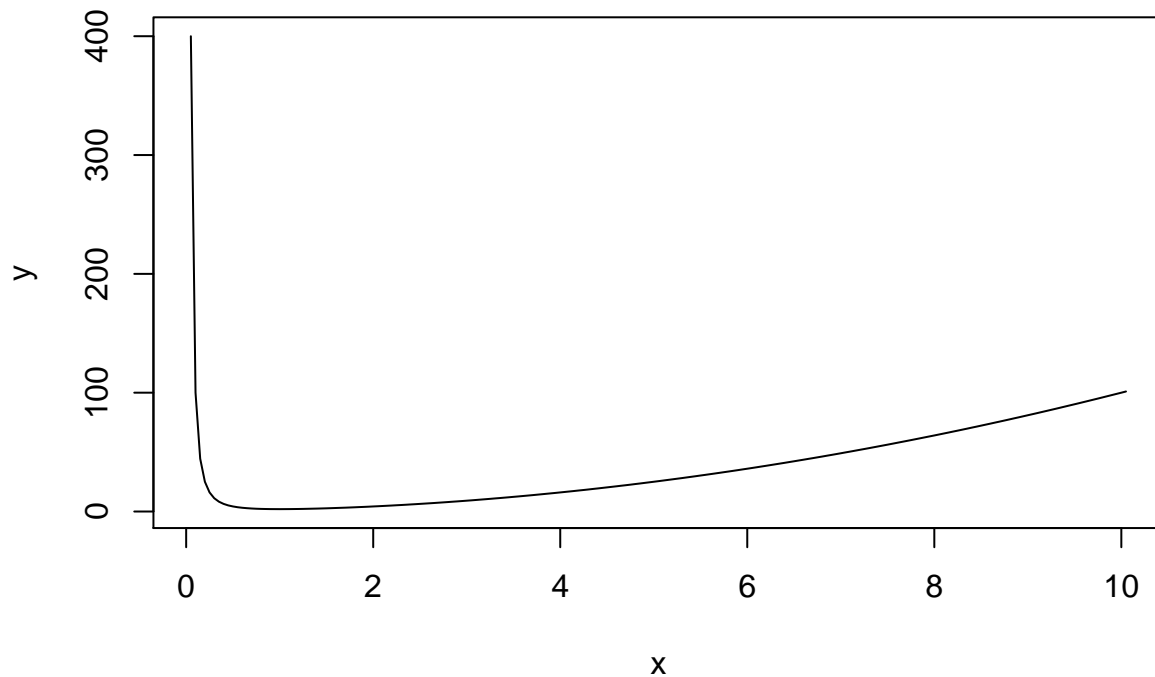
```

g2 <- (-alpha)*2*x/(r2*r2)
g<-as.double(g1+g2)
return(g)
}

x <- seq(-100:100)/20.0
y <- x

for (ii in 1:length(x)){
  y[ii] <- cstick.f(x[ii])
}
plot(x, y, type='l') # ?? does not plot from console??

```



```

x0 <- c(1,2)
require(optimx)

```

```
## Loading required package: optimx
```

```
sdef0 <- optimx(x0, cstick.f, cstick.g, method="Rvmmin", control=list(trace=1))
```

```

## Parameter scaling:[1] 1 1
## gradient test tolerance = 6.055454e-06 fval= 5.2
## compare to max(abs(gn-ga))/(1+abs(fval)) = 6.066115e-12
## Rvmminu -- J C Nash 2009-2015 - an R implementation of Alg 21
## Problem of size n= 2 Dot arguments:
## list()
## Initial fn= 5.2
## ig= 1 gnorm= 4.293251 1 1 5.2
## ig= 2 gnorm= 3.884638 2 2 4.468295
## *ig= 3 gnorm= 2.852834 4 3 3.087837
## *ig= 4 gnorm= 2.154552 6 4 2.520416
## ig= 5 gnorm= 3.731383 7 5 2.417429
## ig= 6 gnorm= 1.086155 8 6 2.098346

```

```
## ig= 7   gnorm= 0.504457      9   7   2.018112
## ig= 8   gnorm= 0.1371067     10   8   2.001136
## ig= 9   gnorm= 0.01377774    11   9   2.000012
## ig= 10  gnorm= 0.0003470064   12  10   2
## ig= 11  gnorm= 8.985444e-07   13  11   2
## ig= 12  gnorm= 5.846616e-11   14  12   2
## *****No acceptable point
## Converged
## Seem to be done Rvmmminu
```

```
sdef0
```

```
## $par
## [1] -0.4472136 -0.8944272
##
## $value
## [1] 2
##
## $counts
## function gradient
##      19      12
##
## $convergence
## [1] 0
##
## $message
## [1] "Rvmmminu appears to have converged"
```

```
xstar <- sdef0$par
gstar <- cstick.g(xstar)
cat("Gradient at proposed solution:")
```

```
## Gradient at proposed solution:
```

```
print(gstar)
```

```
## [1] -2.614686e-11 -5.229372e-11
```

```
## FIXED??
## This doesn't seem to work well??
require(autodiffR)
ad_setup()
hc <- autodiffR::makeHessianFunc(cstick.f)
hstar<-hc(xstar)
cat("Hessian at proposed solution:\n")
```

```
## Hessian at proposed solution:
```

```
print(hstar)
```

```
##      [,1] [,2]
## [1,]  1.6  3.2
## [2,]  3.2  6.4
```

```
print(eigen(hstar)$values)
```

```
## [1] 8.000000e+00 5.846612e-11
```



```

## ?? doesn't seem right
hc(x0)

##      [,1] [,2]
## [1,] 1.984 0.128
## [2,] 0.128 2.176

require(numDeriv)
hcn0 <- numDeriv::hessian(cstick.f, x0)
hcn0

##      [,1] [,2]
## [1,] 1.984 0.128
## [2,] 0.128 2.176

hcnstar <- numDeriv::hessian(cstick.f, xstar)
hcnstar

##      [,1] [,2]
## [1,] 1.6 3.2
## [2,] 3.2 6.4

hcnj0 <- numDeriv::jacobian(cstick.g, x0)
hcnj0

##      [,1] [,2]
## [1,] 1.984 0.128
## [2,] 0.128 2.176

hcnjstar <- numDeriv::jacobian(cstick.g, xstar)
hcnjstar

##      [,1] [,2]
## [1,] 1.6 3.2
## [2,] 3.2 6.4

eigen(hcnstar)$values

## [1] 8.000000e+00 1.137876e-10

```

## Test problem – Wood 4 parameter function

This is reported by Moré, Garbow, and Hillstom (1980) as coming from Colville (1968). The problem in 4 parameters seems to have a false solution far from the accepted one. Is there a good description of this function and the issues it presents?

```

require(autodiff)
ad_setup() # to ensure it is established
#Example 2: Wood function
#
wood.f <- function(x){
  res <- 100*(x[1]^2-x[2])^2+(1-x[1])^2+90*(x[3]^2-x[4])^2+(1-x[3])^2+
    10.1*((1-x[2])^2+(1-x[4])^2)+19.8*(1-x[2])*(1-x[4])
  return(res)
}
#gradient:
wood.g <- function(x){

```

```

g1 <- 400*x[1]^3-400*x[1]*x[2]+2*x[1]-2
g2 <- -200*x[1]^2+220.2*x[2]+19.8*x[4]-40
g3 <- 360*x[3]^3-360*x[3]*x[4]+2*x[3]-2
g4 <- -180*x[3]^2+200.2*x[4]+19.8*x[2]-40
return(c(g1,g2,g3,g4))
}
#hessian:
wood.h <- function(x){
  h11 <- 1200*x[1]^2-400*x[2]+2;    h12 <- -400*x[1]; h13 <- h14 <- 0
  h22 <- 220.2; h23 <- 0;          h24 <- 19.8
  h33 <- 1080*x[3]^2-360*x[4]+2;    h34 <- -360*x[3]
  h44 <- 200.2
  H <- matrix(c(h11,h12,h13,h14,h12,h22,h23,h24,
                h13,h23,h33,h34,h14,h24,h34,h44),ncol=4)
  return(H)
}
#####
x0 <- c(-3,-1,-3,-1) # Wood standard start

cat("Function value at x0=",wood.f(x0),"\n")

## Function value at x0= 19192

wood.ag <- autodiffr::makeGradFunc(wood.f)
cat("Autodiffr gradient value:")

## Autodiffr gradient value:
vwag0<-wood.ag(x0)
print(vwag0)

## [1] -12008 -2080 -10808 -1880
cat("Manually coded:")

## Manually coded:
vwg0 <- wood.g(x0)
print(vwg0)

## [1] -12008 -2080 -10808 -1880
cat("Differences:\n")

## Differences:
print(vwag0-vwg0)

## [1] 0 0 0 0
cat("Autodiffr hessian of function value:")

## Autodiffr hessian of function value:
wood.ah <- autodiffr::makeHessianFunc(wood.f)
vwah0 <- wood.ah(x0)
print(vwah0)

##          [,1] [,2] [,3] [,4]
## [1,] 11202 1200.0      0  0.0

```

```
## [2,] 1200 220.2 0 19.8
## [3,] 0 0.0 10082 1080.0
## [4,] 0 19.8 1080 200.2

cat("Autodiffr hessian via jacobian of autodiff gradient value:")
```

```
## Autodiffr hessian via jacobian of autodiff gradient value:
```

```
wood.ahjag <- autodiffr::makeJacobianFunc(wood.ag)
vwahjag0<-wood.ahjag(x0)
print(vwahjag0)
```

```
##      [,1] [,2] [,3] [,4]
## [1,] 11202 1200.0 0 0.0
## [2,] 1200 220.2 0 19.8
## [3,] 0 0.0 10082 1080.0
## [4,] 0 19.8 1080 200.2
```

```
cat("Autodiffr hessian via jacobian of manual gradient value:")
```

```
## Autodiffr hessian via jacobian of manual gradient value:
```

```
wood.ahj <- autodiffr::makeJacobianFunc(wood.g)
vwahj0 <- wood.ah(x0)
print(vwahj0)
```

```
##      [,1] [,2] [,3] [,4]
## [1,] 11202 1200.0 0 0.0
## [2,] 1200 220.2 0 19.8
## [3,] 0 0.0 10082 1080.0
## [4,] 0 19.8 1080 200.2
```

```
cat("Manually coded:")
```

```
## Manually coded:
```

```
vwh0<-wood.h(x0)
print(vwh0)
```

```
##      [,1] [,2] [,3] [,4]
## [1,] 11202 1200.0 0 0.0
## [2,] 1200 220.2 0 19.8
## [3,] 0 0.0 10082 1080.0
## [4,] 0 19.8 1080 200.2
```

```
cat("Differences from vwh0\n")
```

```
## Differences from vwh0
```

```
cat("vwah0\n")
```

```
## vwah0
```

```
print(vwah0-vwh0)
```

```
##      [,1] [,2] [,3] [,4]
## [1,] 0 0 0 0
## [2,] 0 0 0 0
## [3,] 0 0 0 0
## [4,] 0 0 0 0
```

```

cat("\n")

cat("vwahj0\n")

## vwahj0
print(vwahj0-vwh0)

##      [,1] [,2] [,3] [,4]
## [1,]    0    0    0    0
## [2,]    0    0    0    0
## [3,]    0    0    0    0
## [4,]    0    0    0    0

cat("\n")

cat("vwahjag0\n")

## vwahjag0
print(vwahjag0-vwh0)

##      [,1] [,2] [,3] [,4]
## [1,]    0    0    0    0
## [2,]    0    0    0    0
## [3,]    0    0    0    0
## [4,]    0    0    0    0

cat("\n")

## d <- c(1,1,1,1)
require(optimx)
meths <- c("snewton", "snewtonm", "nlm")
wdefault <- opm(x0, fn=wood.f, gr=wood.g, hess=wood.h, method=meths, control=list(trace=0))
print(wdefault)

##      p1 p2 p3 p4      value fevals gevals convergence kkt1 kkt2
## snewton  1  1  1  1 1.142616e-29   119    70           92 TRUE TRUE
## snewtonm  1  1  1  1 1.399599e-28    88    50            0 TRUE TRUE
## nlm      1  1  1  1 1.004943e-16    NA   335            0 TRUE TRUE
##      xtime
## snewton  0.008
## snewtonm 0.004
## nlm      0.004

wagah <- opm(x0, fn=wood.f, gr=wood.ag, hess=wood.ah, method=meths, control=list(trace=0))

## Small gradient
print(wagah)

##      p1 p2 p3 p4      value fevals gevals convergence kkt1 kkt2
## snewton  1  1  1  1 0.000000e+00   116    67            0 TRUE TRUE
## snewtonm  1  1  1  1 0.000000e+00    81    50            0 TRUE TRUE
## nlm      1  1  1  1 1.004942e-16    NA   335            0 TRUE TRUE
##      xtime
## snewton  3.868
## snewtonm 2.912
## nlm      22.724

```

```
## Timings

thand <- microbenchmark(wood.h(x0))
tad <- microbenchmark(wood.ah(x0))
print(thand)

## Unit: microseconds
##      expr    min      lq    mean median      uq    max neval
## wood.h(x0) 4.806 5.052 5.54656 5.2405 5.612 20.994   100

print(tad)

## Unit: milliseconds
##      expr    min      lq    mean  median      uq    max neval
## wood.ah(x0) 27.57984 28.88468 30.51616 29.42614 30.06532 59.8322   100
```

## Performance issues

Optimization is, by its very nature, about improving things. Thus it is of prime interest to seek faster and better ways to optimize functions. In this section we look at some issues that may influence the speed, reliability and correctness of optimization calculations.

First, it is critical to note that **R** almost always offers several ways to accomplish the same computational result. However, the speed with which the different approaches return a result can be wildly different. (?? can JN find the 800% scale factor??).

Second, there are many parts of the autodiff wrapper of Julia's automatic differentiation that may use up computing cycles:

- We must translate from one programming language to another in some sense in order to call the appropriate functions in Julia based on **R** functions.
- Results must be properly structured on return to **R**.
- Hand coded derivative expressions, especially hand-optimized ones, can be expected to out-perform automatic differentiation results.

NOTE: Performance is interesting, but it is far from the complete picture. We can use results from autodiff to validate hand-coded functions. We can get results that are efficient of human time and effort that may be otherwise unavailable. Moreover, the results of computing gradients and Hessians allow us to conclude that a solution has been achieved.

### A small performance comparison using autodiff

```
rm(list=ls())
require(autodiff)
autodiff::ad_setup() # to ensure it is established

ores <- function(x){
  x # Function will be the parameters. ofn is sum of squares
}

logit <- function(x) exp(x) / (1 + exp(x))

ofn <- function(x){
```

```

    res <- ores(x) # returns a vector of residual values
    sum(logit(res) ^ 2)
}

## Now try to generate the gradient function
ogr <- autodiffr::makeGradFunc(ofn)

system.time(ogr(runif(100)))

##      user  system elapsed
##    0.280    0.000    0.281
system.time(ogr(runif(100)))

##      user  system elapsed
##    0.016    0.000    0.015
ogr1 <- autodiffr::makeGradFunc(ofn, x = runif(100))

system.time(ogr1(runif(100)))

##      user  system elapsed
##    0.008    0.000    0.007
system.time(ogr1(runif(100)))

##      user  system elapsed
##    0.008    0.000    0.008
ogr2 <- autodiffr::makeGradFunc(ofn, x = runif(100), use_tape = TRUE)

system.time(ogr2(runif(100)))

##      user  system elapsed
##    0.140    0.000    0.142
system.time(ogr2(runif(100)))

##      user  system elapsed
##    0.000    0.000    0.001

```

## A problem with discontinuous gradient

Problems with discontinuous gradient may give gradient methods difficulty.

Here is a problem where the gradient is discontinuous, but not at the minimum.

*## discontin.R, a test function with discontinuous gradient*

```

disc.f <- function(x){
  nn <- length(x)
  val <- 0.0
  for (ii in 1:nn){
    tt <- (x[ii] - ii)
    if (abs(tt) < ii) {
      ff <- tt*tt
    } else {

```

```

      ff <- abs(tt)
    }
    val <- val + ff*ff
  }
  val
}

require(optimx)
x0 <- runif(4)
x0

## [1] 0.89078635 0.47916051 0.94314744 0.05917632

sol0 <- optimr(x0, disc.f, method="nmkb")

## Warning in optimr(x0, disc.f, method = "nmkb"): Successful convergence
## Restarts for stagnation =0

sol0

## $par
## [1] 0.9978913 2.0128136 3.0128012 3.9996128
##
## $value
## [1] 5.383122e-08
##
## $convergence
## [1] 0
##
## $message
## [1] "Successful convergence"
##
## $counts
## [1] 85 NA
##
## $nitns
## [1] NA

require(autodiffr)
ad_setup()
disc.ag <- autodiffr::makeGradFunc(disc.f)
sol1 <- optimr(x0, disc.f, disc.ag, method="Rvmmin", control=list(trace=1))

## Parameter scaling:[1] 1 1 1 1
## gradient test tolerance = 6.055454e-06 fval= 264.432
## compare to max(abs(gn-ga))/(1+abs(fval)) = 3.777694e-11
## Rvmminu -- J C Nash 2009-2015 - an R implementation of Alg 21
## Problem of size n= 4 Dot arguments:
## list()
## Initial fn= 264.432
## ig= 1 gnorm= 247.6675 1 1 264.432
## **ig= 2 gnorm= 12.27613 4 2 35.27634
## *ig= 3 gnorm= 11.32922 6 3 32.08908
## ig= 4 gnorm= 19.51258 7 4 30.11557
## *ig= 5 gnorm= 13.34165 9 5 24.49628
## ig= 6 gnorm= 33.33428 UPDATE NOT POSSIBLE: ilast, ig 1 6

```

```

## 10 6 20.04132
## *ig= 7 gnorm= 8.842687 12 7 18.74412
## **ig= 8 gnorm= 8.730798 15 8 18.2288
## ig= 9 gnorm= 4.487915 16 9 5.032437
## ig= 10 gnorm= 0.8707973 17 10 0.1467999
## ig= 11 gnorm= 0.05637104 18 11 0.00405107
## ig= 12 gnorm= 0.03706002 19 12 0.002336852
## ig= 13 gnorm= 0.01348833 20 13 0.0006334487
## ig= 14 gnorm= 0.00634206 21 14 0.0002456539
## ig= 15 gnorm= 0.002918795 22 15 9.445441e-05
## ig= 16 gnorm= 0.001628752 23 16 4.428581e-05
## ig= 17 gnorm= 0.001110674 24 17 2.472968e-05
## ig= 18 gnorm= 0.0008537438 25 18 1.603547e-05
## ig= 19 gnorm= 0.0006529062 26 19 1.087156e-05
## ig= 20 gnorm= 0.0004257771 27 20 6.226156e-06
## *ig= 21 gnorm= 0.000370275 29 21 5.295728e-06
## ig= 22 gnorm= 0.0003074603 30 22 5.009866e-06
## ig= 23 gnorm= 0.0003117494 31 23 4.579412e-06
## ig= 24 gnorm= 0.0003033878 32 24 4.476951e-06
## ig= 25 gnorm= 0.0002469098 33 25 3.687131e-06
## ig= 26 gnorm= 0.0001985193 34 26 2.806005e-06
## ig= 27 gnorm= 0.0001028793 35 27 1.1366e-06
## ig= 28 gnorm= 4.893569e-05 36 28 4.108249e-07
## ig= 29 gnorm= 2.247952e-05 37 29 1.434972e-07
## ig= 30 gnorm= 1.02367e-05 38 30 5.025503e-08
## ig= 31 gnorm= 4.538228e-06 39 31 1.706236e-08
## ig= 32 gnorm= 2.005941e-06 40 32 5.851909e-09
## ig= 33 gnorm= 1.152225e-06 41 33 2.972691e-09
## ig= 34 gnorm= 1.101338e-06 42 34 2.637216e-09
## ig= 35 gnorm= 1.110738e-06 43 35 2.605406e-09
## ig= 36 gnorm= 1.117989e-06 44 36 2.543172e-09
## ig= 37 gnorm= 1.109585e-06 45 37 2.401539e-09
## ig= 38 gnorm= 1.050001e-06 46 38 2.151045e-09
## ig= 39 gnorm= 9.52997e-07 47 39 1.899759e-09
## ig= 40 gnorm= 9.23711e-07 48 40 1.826922e-09
## ig= 41 gnorm= 9.228179e-07 49 41 1.821823e-09
## ig= 42 gnorm= 9.214369e-07 50 42 1.819489e-09
## ig= 43 gnorm= 9.168678e-07 51 43 1.815625e-09
## ig= 44 gnorm= 9.132603e-07 52 44 1.813718e-09
## ig= 45 gnorm= 9.086423e-07 53 45 1.81074e-09
## ig= 46 gnorm= 9.007073e-07 54 46 1.803488e-09
## ig= 47 gnorm= 8.855003e-07 55 47 1.783983e-09
## ig= 48 gnorm= 8.553805e-07 56 48 1.730943e-09
## ig= 49 gnorm= 7.95533e-07 57 49 1.589628e-09
## ig= 50 gnorm= 6.782673e-07 58 50 1.258063e-09
## ig= 51 gnorm= 4.720246e-07 59 51 7.318557e-10
## ig= 52 gnorm= 2.577747e-07 60 52 3.149334e-10
## ig= 53 gnorm= 1.258113e-07 61 53 1.205848e-10
## ig= 54 gnorm= 5.673724e-08 62 54 4.223402e-11
## ig= 55 gnorm= 2.62043e-08 63 55 1.670769e-11
## ig= 56 gnorm= 2.090967e-08 64 56 1.296368e-11
## ig= 57 gnorm= 1.971819e-08 65 57 1.179054e-11
## ig= 58 gnorm= 1.439377e-08 66 58 7.173373e-12
## ig= 59 gnorm= 9.262742e-09 67 59 4.156835e-12

```



## ig= 60	gnorm= 6.942612e-09	68	60	2.985614e-12
## ig= 61	gnorm= 5.402857e-09	69	61	2.161985e-12
## ig= 62	gnorm= 4.447603e-09	70	62	1.686036e-12
## ig= 63	gnorm= 4.141578e-09	71	63	1.516493e-12
## ig= 64	gnorm= 3.97784e-09	72	64	1.409317e-12
## ig= 65	gnorm= 3.844147e-09	73	65	1.309976e-12
## ig= 66	gnorm= 3.775019e-09	74	66	1.229767e-12
## ig= 67	gnorm= 3.77391e-09	75	67	1.200132e-12
## ig= 68	gnorm= 3.758109e-09	76	68	1.190344e-12
## ig= 69	gnorm= 3.692948e-09	77	69	1.142225e-12
## ig= 70	gnorm= 3.624243e-09	78	70	1.105634e-12
## ig= 71	gnorm= 3.360546e-09	79	71	1.003201e-12
## *ig= 72	gnorm= 3.279417e-09	81	72	9.82674e-13
## ig= 73	gnorm= 3.149905e-09	82	73	9.643588e-13
## ig= 74	gnorm= 3.215072e-09	83	74	9.521131e-13
## ig= 75	gnorm= 3.203421e-09	84	75	9.502317e-13
## ig= 76	gnorm= 3.198528e-09	85	76	9.500354e-13
## ig= 77	gnorm= 3.199201e-09	86	77	9.500134e-13
## ig= 78	gnorm= 3.200698e-09	87	78	9.499855e-13
## ig= 79	gnorm= 3.201162e-09	88	79	9.499704e-13
## ig= 80	gnorm= 3.203005e-09	89	80	9.498732e-13
## ig= 81	gnorm= 3.205116e-09	90	81	9.496754e-13
## ig= 82	gnorm= 3.208405e-09	91	82	9.491043e-13
## ig= 83	gnorm= 3.212159e-09	92	83	9.476752e-13
## ig= 84	gnorm= 3.214528e-09	93	84	9.439255e-13
## ig= 85	gnorm= 3.207566e-09	94	85	9.342785e-13
## ig= 86	gnorm= 3.163386e-09	95	86	9.081873e-13
## ig= 87	gnorm= 2.966203e-09	96	87	8.430857e-13
## ig= 88	gnorm= 2.919437e-09	97	88	8.111538e-13
## ig= 89	gnorm= 2.635285e-09	98	89	7.132601e-13
## ig= 90	gnorm= 1.055377e-09	99	90	2.161785e-13
## ig= 91	gnorm= 5.135445e-10	100	91	8.547025e-14
## ig= 92	gnorm= 2.692084e-10	101	92	3.867616e-14
## ig= 93	gnorm= 1.238153e-10	102	93	1.375787e-14
## ig= 94	gnorm= 5.282327e-11	103	94	4.377101e-15
## ig= 95	gnorm= 2.335674e-11	104	95	1.452558e-15
## ig= 96	gnorm= 1.031369e-11	105	96	4.774858e-16
## ig= 97	gnorm= 4.63775e-12	106	97	1.593293e-16
## ig= 98	gnorm= 2.119295e-12	107	98	5.378859e-17
## ig= 99	gnorm= 9.873113e-13	108	99	1.851187e-17
## ig= 100	gnorm= 4.661807e-13	109	100	6.49268e-18
## ig= 101	gnorm= 2.203452e-13	110	101	2.299234e-18
## ig= 102	gnorm= 1.025669e-13	111	102	8.040469e-19
## ig= 103	gnorm= 4.717145e-14	112	103	2.850656e-19
## ig= 104	gnorm= 2.490176e-14	113	104	1.413987e-19
## ig= 105	gnorm= 2.060521e-14	114	105	1.228037e-19
## ig= 106	gnorm= 2.039198e-14	115	106	1.223164e-19
## ig= 107	gnorm= 2.037204e-14	116	107	1.222475e-19
## ig= 108	gnorm= 2.023421e-14	117	108	1.215244e-19
## ig= 109	gnorm= 2.007859e-14	118	109	1.201898e-19
## ig= 110	gnorm= 1.979703e-14	119	110	1.16419e-19
## ig= 111	gnorm= 1.935424e-14	120	111	1.083769e-19
## ig= 112	gnorm= 1.834558e-14	121	112	9.313666e-20
## ig= 113	gnorm= 1.585084e-14	122	113	7.129666e-20

```
## ig= 114    gnorm= 1.13452e-14    123    114    4.520048e-20
## ig= 115    gnorm= 6.341097e-15    Seem to be done Rvmmminu
```

```
soll
```

```
## $par
## [1] 1.000005 1.999992 3.000011 3.999991
##
## $value
## [1] 2.536705e-20
##
## $counts
## function gradient
##      124      115
##
## $convergence
## [1] 2
##
## $message
## [1] "Rvmmminu appears to have converged"
```

?? Do we want to try discontinuity at solution? Discontinuous function value?

## Bibliography

- Colville, A R. 1968. "A Comparative Study of Nonlinear Programming Codes, Rep. 320-2949." IBM New York Scientific Center.
- Fletcher, R. 1965. "Function Minimization Without Calculating Derivatives – a Review." *Computer Journal* 8: 33–41.
- Moré, J. J., B. S. Garbow, and K. E. Hillstom. 1980. "ANL-80-74, User Guide for MINPACK-1." Argonne National Laboratory. <http://www.mcs.anl.gov/~jmore/ANL8074a.pdf>.
- Nash, John C. 1979. *Compact Numerical Methods for Computers: Linear Algebra and Function Minimisation*. Bristol: Adam Hilger.